

Investor Persona: Tier 1 VC

Synthesised from 20VC, YC W26, Sequoia, a16z, Index, FT, and X content from Nov 2025 – May 2026. Treat as a composite — a portrait of the consensus mind among partners at firms like Sequoia, a16z, Index, Founders Fund, Greylock, who read this material and update from it weekly.

Chapter 1 — Snapshot

He has spent the last six months watching the centre of gravity in venture move twice. First, away from "AI as a feature" toward agents that **actually do work**. Second, away from the assumption that all of this would express itself in better SaaS and toward the assumption that it expresses itself in **services revenue, physical-world systems, and infrastructure**. His most-read piece this year is probably Sequoia's 2026: *This is AGI*. His most-cited data point is probably that YC W26 had **14 companies at \$1M ARR by Demo Day** — a record — with **3x more than W25** and a batch-wide average of **14% week-on-week growth**.

He is operating in a market that is simultaneously the most capital-rich and the most punishing in recent memory. **Q1 2026 saw roughly \$297B of global VC deployed, but three deals — OpenAI, Anthropic, xAI — accounted for ~67% of it**. He knows what this means: the middle has collapsed. He is either writing into the top of the market at prices that look insane unless you believe in trillion-dollar outcomes, or he is hunting at seed/Series A for companies that look nothing like the 2021 vintage.

He doesn't want to hear "AI-powered." He wants to hear what work you complete, what budget you replace, what physical or regulated constraint you own that a foundation model cannot eat, and why your company gets **stronger** as models get better — not survivable.

Chapter 2 — Thesis Areas (Where He's Leaning In)

2.1 Long-Horizon Agents — "Doers, Not Talkers"

This is the single most important shift in his worldview from late 2025 to now. Sequoia's January 2026 essay argues that long-horizon agents are functionally AGI — the third inflection point after pre-training (2022) and inference-time compute (late 2024). The key claim he has internalised: the rate of progress is exponential, with agent task-completion duration doubling roughly every seven months — agents able to do a full day of expert work by 2028, a full year by 2034.

His pitch-meeting filter is now: "Can you hire it?" Sarah Guo's litmus test, repeated approvingly by Sequoia, has become the standard partner-meeting question. The applications of 2023–2024 were talkers; the applications of 2026–2027 will be doers — that's the line you'll hear him use.

What he wants to see in a pitch:

- An agent that **persists, recovers from failure, and finishes the job** — not one that returns a draft for a human to fix
- Evidence of an **agent harness** (memory, evals, tool use, error recovery) that you've built around the foundation model — not just a system prompt
- A specific human role you replace, with the **billable hours/headcount it eliminates**
- Reliability metrics. He has heard a hundred demos. He wants pass rates on real workloads, not vibes.

2.2 AI-Native Services / "Selling Work"

This is the commercial expression of the agent thesis and the single biggest TAM expansion narrative he's pitching to his own LPs. Sequoia partners argue the prize is roughly \$10 trillion in services revenue the software industry has never been able to touch — US legal services alone is a \$400B market, equivalent to all global software spend.

He has watched the YC batches re-shape themselves around this. AI-native services — where AI performs the job end-to-end with humans supervising — was the largest software category in W26 at 28%. Companies have moved from building tools for regulated industries to running the professional services themselves: General Legal as an AI-native law firm charging flat \$500/contract fees, LegalOS running immigration cases at 99%+ visa approval, Beacon Health deploying AI employees for primary care. He thinks of this as **capturing professional services margins, not SaaS margins**.

What gets him to write a cheque:

- You own a **license, regulatory trust, or professional credential** that lets you actually deliver the service
- You price on **outcomes** (per contract, per case, per claim) — not per seat
- You can show **margin structure that approaches SaaS over time** as the AI does more of the work
- You're attacking a vertical where the **existing labour cost is mostly cognitive and mostly billable**

2.3 American Dynamism / The AI-Native Industrial Base

If you're pitching anything physical-world, this is the lens he uses. a16z's Big Ideas 2026 frames it as the rebuilding of energy, manufacturing, logistics, and infrastructure — an industrial base that is AI-native and software-first from day one, starting with simulation, automated design, and AI-driven operations. The opportunities: advanced energy systems, robotics heavy manufacturing, next-generation mining, biological and enzymatic precursor chemicals, autonomous fleet coordination.

The geopolitical backdrop matters to him. The US once led in 60 of the 64 most critical technologies; China now leads in 57 and the US in only 7. He sees American Dynamism as a

generational rebuild, not a sector bet.

Specific sub-themes you should know he cares about:

- Physical observability — autonomous sensors and drones giving continuous visibility into ports, rail, power lines, pipelines, military bases, datacenters
- Autonomous labs that close the loop on scientific discovery — hypothesis to experiment to iteration without human intervention
- The "electro-industrial stack" — batteries, motors, power electronics, software powering EVs and autonomous machines
- Data crusade in critical industries — capturing the unstructured operational data sitting in energy, manufacturing, logistics

What he wants to see:

- A factory-style approach to a problem people assume is bespoke (he's heard the Musk-ism "the factory is the product" a hundred times and still likes it)
- A founder who can talk credibly about both the bits *and* the atoms
- A path to government or critical-infrastructure procurement — he no longer treats DoD as a dirty word

2.4 Deep Tech & Hardware

The W26 batch broke his mental model of YC. One in eight W26 companies is building something physical — robots, drones, wearables, space hardware, biotech. A 22-year-old dropout building the first hotel on the Moon presented a moon brick to Congress. YC's Spring 2026 RFS dedicates three of ten themes to companies operating in atoms: "Modern Metal Mills," "Large Spatial Models," "AI Guidance for Physical Work". He reads this as YC putting institutional weight behind a trend that was already showing up in his deal flow.

He also believes there's froth here. From the Demo Day investor write-ups: the highest valuations in W26 were not the high-revenue companies — they were the deep tech companies, suggesting a bubble driven by fear that foundation models will eat application companies. He'll fund this, but he'll do diligence harder than founders expect.

What he's looking for:

- A **proprietary physical-world constraint** that can't be replicated with software alone — manufacturing know-how, proprietary materials, hard-won regulatory clearance, a sensor network
- Founders with **real domain depth** — not "AI for X" wrappers
- A credible path to **defence, energy, or critical-infrastructure revenue** within 24 months

2.5 Agentic Infrastructure

His "picks-and-shovels" bet. By Winter 2026 the funding moved from "AI agent for Y

industry" to building the plumbing underneath all of those agents — auth, testing, security, context, observability, billing rails. He likes the historical parallel: infrastructure companies have ended up worth more than application companies in every previous platform shift — Stripe vs any e-commerce store, Datadog vs the apps it monitors.

He knows the window is short. The window to become the default tool in your layer is maybe 12-18 months before consolidation hits and late entrants get squeezed. He also knows the hyperscaler risk is real: AWS / GCP / Azure are building competing tooling internally.

What he wants:

- A specific layer (auth for agents, eval, memory, browser automation, sandboxing, payment rails) where you can be the default
- **Developer love** as the leading indicator — GitHub stars, real deploys, design-partner logos
- A reason you survive the hyperscaler-launches-a-competitor moment (community, switching costs, depth in one workflow)

2.6 The Innermost Loop — Compute, Energy, Tokens

His macro frame for the entire AI cycle. OpenAI's Greg Brockman has publicly said the company faces an acute compute shortage despite aggressive procurement, with GPU availability in 2026 effectively rounding to zero. The five largest US hyperscalers — Microsoft, Alphabet, Amazon, Meta, Oracle — have collectively committed to \$660-690 billion of capex in 2026, nearly doubling 2025. Goldman estimates as much as \$200B of upside to current hyperscaler 2026 capex estimates, with 2027 likely above \$1 trillion.

His operating belief: **compute is the binding constraint on everything else**. That makes everything that produces, moves, cools, or schedules compute interesting — including some categories that would have been laughed at two years ago (orbital data centres, behind-the-meter power, novel cooling, custom silicon).

What he wants to see:

- A real technical edge — not just "we'll build a data centre too"
- A demand-side anchor (a hyperscaler commitment, a sovereign customer, a frontier lab LOI)
- A capital plan that doesn't require him personally to fund the whole build

Chapter 3 — Anti-Thesis (What He's Cooled On)

3.1 Seat-Based SaaS — The "SaaS Massacre" of Early 2026

He no longer believes the median seat-based SaaS company is a venture-scale business. The 20VC weekly debate explicitly called the early-2026 "SaaS massacre" and reset the IPO

threshold at \$4B revenue plus 50% growth, with 10-15x FCF as the floor for ex-growth SaaS. Public SaaS now trades at 8-9x cash flow. The bifurcation in his deal flow mirrors that of the public markets.

If you walk in pitching a per-seat collaboration tool, his energy will drop visibly. He's heard the agent thesis enough times to believe that **headcount-based pricing is the wrong unit for a world where the headcount you serve is shrinking**.

How to handle it if you are in this space:

- Reframe pricing around **outcome, throughput, or work units** — not seats
- Show that your **net dollar retention holds even as customer seat counts decline**
- Have a clear story for **how AI makes you stronger, not survivable**

3.2 "Copilots" and AI Wrappers

The word "copilot" has gone from selling point to red flag in his vocabulary. The share of YC companies calling themselves a "copilot" dropped from about 4% in early 2025 to 1% in W26 — copilots assist, agents act, and the industry moved on. He has personally listened to multiple Sequoia AI Ascent talks reframing the gap. Pat Grady's "faster horses vs cars" framing — applications that make you 10-40% more productive don't change the nature of work, applications that make you 10-40x more productive do — has become his shorthand.

Translation: he doesn't want to fund anything where the customer still does most of the work and your product just makes it nicer.

3.3 Application-Layer Point Solutions

He worries about feature-risk from foundation labs. He'll bring up the "what happens when GPT-5.5 / Claude 5 / Gemini 3 ships this natively" objection in the first 15 minutes of the meeting. The bar to clear this objection has risen: you need **workflow depth, proprietary data, distribution, or a non-software moat** — not just a better UI.

This is also why he's gravitated toward services (you own the customer relationship and the licence) and infrastructure (you're below the foundation model, not on top of it).

3.4 Consumer (with Narrow Exceptions)

W26 was 64% B2B with only ~5% consumer. Garry Tan's published data showed that for investors who backed at least three companies per Demo Day from 2018-2020, the bottom quartile of Demo Day investing outperformed the top quartile of the broader VC market — and yet consumer keeps shrinking as a share. He's pattern-matched on this.

He'll still fund consumer if it's a **new behaviour enabled by AI** (e.g., the AI sandbox game, the AI wearable that hit \$27M ARR with 50% MoM growth), but the default answer is no.

3.5 The Mid-Growth Middle

This is the most important anti-thesis to understand because it's structural, not stylistic. Index's Mike Volpi said the venture market has bifurcated — a handful of companies

raising rounds that look like sovereign debt issuances, and everyone else competing for a shrinking pool. The middle has been hollowed out.

He's not going to write a \$30M Series B into a company growing 60% YoY at \$15M ARR with no obvious path to either acquisition or IPO. That cohort got buried in 2025-26. He'll either go earlier (where prices are sane and outcomes can be 100x) or later (where the trillion-dollar-outcome math works).

What this means for you: if you are at \$5-25M ARR growing 60-100%, you need to either accelerate to triple-triple territory or accept a smaller round at a flatter price — and don't pitch him expecting a 2021-vintage Series B markup.

Chapter 4 — Mental Models & Frameworks He Uses

These are the lenses he applies in partner meetings. Mirror them in your deck and he'll feel you're operating at his altitude.

Talkers vs Doers. The Sequoia frame. Talkers respond; doers persist, recover, and complete work. Pitch the work completed, not the conversation enabled.

AI Velocity / Time-to-Revenue Compression. Hex Security hit \$1M ARR in eight weeks; Luel approached \$2M ARR in six weeks. AI has compressed the timeline from idea to \$1M ARR by roughly 10x versus the traditional 12-18 months for SaaS. He uses this as both a positive selector (top companies move this fast) and a negative one (if you've been at it 18 months without traction, the AI tailwind didn't help you).

The Liquidity Paradox / Bifurcation. Record capital alongside record concentration. He thinks about every deal in this context — either you're a megadeal candidate (trillion-dollar-outcome math) or you're competing for the residual pool with much sharper terms.

MAD Framework (Sequoia). Mode (which model and harness), Affordances (what unique workflow/data you can act on), Distribution. He'll probe all three.

Fat Startups Win. The 2010 Horowitz-vs-Wilson "fat vs lean" debate is now considered settled in favour of fat — AI-era startups should build the full stack, own distribution, ship outcomes not features.

Trillion-Dollar Outcome Math. The 20VC framing: the first question in your investment memo should be "odds of trillion-dollar outcome." If it's north of 2%, do the deal. This is why he'll write into Perplexity at \$14B or Anthropic at \$380B — he's not buying the multiple, he's buying the at-bat.

Mark-to-Reality. Rory O'Driscoll's heuristic — take a 2021 unicorn valuation and run it through tiers asking what's the real growth net of price increases, what tier does it belong to, what's implied fair value today. The answers are often uncomfortable. He applies this to every late-stage deal that walks in.

Service-as-Software. The model where AI does 80% of the work and humans provide oversight and licences. The company captures professional services margins rather than SaaS tool margins.

The Innermost Loop. Capital is shifting to the chips, energy, and tokens needed to power the agentic economy. He uses this to rationalise hardware/energy bets that would have looked exotic two years ago.

Chapter 5 — The New Investment Bar

He has been recalibrated by W26. The numbers that now define "elite" in his head:

- **\$1M ARR by Demo Day** for application companies. 14 of ~200 W26 companies hit this milestone, 7% of the batch versus the historical 2-3%.
- **14% week-on-week revenue growth** as the batch average — the highest growth rate ever recorded for a YC cohort.
- **Lean teams.** 22 solo founders in W26 — 11% of the batch is one person building alone. Average founder has 5.8 years of professional experience, down from the historical ~9; AI agent founders sit at 4.8 years.
- **AI-native from day one.** Not retrofitted.
- **A proprietary non-software constraint.** The W26 winners own something that can't be replicated with software: physical-world data, manufacturing expertise, regulatory trust, or a professional licence.
- **Valuation reality.** \$1M-ARR W26 companies typically priced at \$25-40M post; deep-tech companies pushed to \$50-100M with some outliers at \$200M. He'll pay the lower end, hesitate at the top.

If you're walking in below these benchmarks at seed/Series A, you need a very specific reason — non-recurring revenue, a hardware build cycle, a regulated sector with longer sales cycles.

Chapter 6 — Vocabulary & Frame

Words to use (he uses these unprompted): doer, agent harness, long-horizon, persistence, recovery, evals, AI-native, service-as-software, selling work, outcome pricing, vertical agent, physical-world data, sovereign demand, electro-industrial stack, atoms, factory mindset, MAD, trillion-dollar at-bat, innermost loop, mark-to-reality.

Words to avoid (these register as 2024 vocabulary): copilot, AI-powered, wrapper, GPT wrapper, productivity tool, seamless, enterprise-grade, democratise, ChatGPT for X, AI assistant, augmented worker.

Framing he respects:

- "We sell completed work, not software seats"
- "Our agent's reliability on the canonical workload is X% — here's the eval set"
- "We get stronger as models get smarter because [proprietary distribution / data / regulatory position]"
- "Our outcome is priced at a fraction of the billable hours we replace"
- "Our customer is the [licensed professional / regulated buyer / DoD programme office]"

Framing he distrusts:

- "We're an AI-first version of [existing SaaS category]"
 - "Our moat is the data flywheel" (he's heard this 500 times — show him the data, not the diagram)
 - "We'll partner with the hyperscalers" (he assumes hyperscalers will eat you)
-

Chapter 7 — How to Pitch Him (Mapped to Your Sequoia Template)

Company purpose. Lead with the work your company completes, not the tech it uses. "We are the AI-native [law firm / radiologist / underwriter] for X." Not "We use AI to help X."

Problem. Quantify it in the right unit for 2026: billable hours, headcount, throughput, error rate. "Mid-market law firms spend \$X on commercial contract review with 48-hour turnaround." Show you understand the actual workflow, not just the category.

Solution. Describe the **agent harness and the failure modes you've engineered around** — not the model you use. Reliability metric on real workloads. The eureka moment should be about a capability that became possible only in the last 6–12 months (otherwise: why hasn't someone done it?).

Why now. This is where you mirror his worldview directly. The three-ingredient stack — pre-training, inference-time compute, long-horizon agents — only crossed the threshold for your workflow recently. Cite specific capability shifts. If you can tie it to Sequoia's exponential, even better.

Market potential. Use the **services TAM** framing, not the software TAM. The right denominator is the labour cost you replace, not a Gartner SaaS market size. If you're in American Dynamism, show the **federal procurement and reshoring tailwinds** — he reads the a16z American Dynamism Summit material and will know the numbers.

Competition / alternatives. Three categories: (1) the existing service providers / labour you're replacing — show economics; (2) other AI-native entrants in your vertical — show why you win; (3) the foundation model labs themselves — show why GPT-5.5 native shipping this doesn't kill you. Skip this third one at your peril.

Business model. Outcome-based pricing wherever possible. Per-case, per-claim, per-contract, per-report. If you must use seats or subscriptions, show why your contract value grows even as the customer's seat count falls.

Team. Lean and technical is the W26 archetype. He's seen solo founders ship \$1M ARR. Don't pad the team slide; show velocity. Domain depth matters more than logos for vertical plays. For American Dynamism / hardware: show people who've built the physical thing before.

Financials. If you have revenue, lead with **week-on-week or month-on-month growth**, not annualised numbers. Show **gross margins trending toward software economics** even if you start service-shaped. Show **capital efficiency** — ARR per FTE is now a benchmark he tracks.

Vision. Five-year vision should sketch the **services-scale TAM** you eventually address (\$X00B labour market, not \$X0B software market). And it should describe what your company looks like when models are 100x more capable — your story should be that you become **more valuable**, not commoditised.

Chapter 8 — Recent Shifts in His Thinking (Nov 2025 → May 2026)

- **Nov-Dec 2025:** a16z publishes Big Ideas 2026 Part 2, formalising American Dynamism as a generational thesis with the "AI-native industrial base" framing. He starts asking every consumer/SaaS pitch why they're not building atoms instead.
- **Jan 2026:** Sequoia publishes "2026: This is AGI" arguing long-horizon agents are functionally AGI. The "talkers vs doers" line becomes ubiquitous in partner meetings.
- **Late Jan - Feb 2026:** The 20VC weekly debate (Stebbins, Lemkin, O'Driscoll) names the "SaaS massacre" — new IPO threshold of \$4B revenue / 50% growth, 10-15x FCF as floor for ex-growth SaaS. He starts revaluing his late-stage SaaS portfolio against this.
- **Q1 2026:** \$255.5B of AI funding globally with OpenAI, Anthropic, xAI accounting for 67% of capital. He internalises the bifurcation — either trillion-dollar at-bats or early-stage with sane prices.
- **March 2026:** YC W26 Demo Day. 14 companies at \$1M ARR, an all-time record. Average 14% WoW growth across the batch. His mental model of "elite seed company" resets.
- **April 2026:** Sequoia AI Ascent. Pat Grady, Sonya Huang, Konstantine Buhler put a \$10 trillion services-revenue prize on the table and project 99.9% of cognitive work eventually automated. The services-as-software thesis crystallises.
- **April 2026:** a16z American Dynamism Summit in DC reinforces the geopolitical urgency — US lost technological lead in 53 of 64 critical categories to China. He sharpens his defence/energy thesis.
- **May 2026:** Hyperscaler 2026 capex revised upward to \$800–900B, 2027 estimates above \$1 trillion. The "compute is the binding constraint" frame becomes his default

Appendix — Sources Behind This Persona

Primary:

- Sequoia: *2026: This is AGI* (Grady, Huang), AI Ascent 2026 keynote, *Services: The New Software*
- a16z: *Big Ideas 2026 Part 2* (American Dynamism + Apps), American Dynamism Summit 2026
- 20VC weekly debate: episodes #18-#45 (Stebbing, Lemkin, O'Driscoll)
- YC: W26 batch breakdown, Spring 2026 RFS, Garry Tan's commentary

Supporting / atmospheric:

- Extruct AI W26 analysis (199 companies, 14 columns)
- The VC Corner W26 database (Rebel Theorem 4.0 scoring)
- PitchBook / Crunchbase / KPMG Q1 2026 venture data
- Goldman Sachs and Morgan Stanley AI capex research
- Index Ventures partner commentary on bifurcation
- Sameer Nanda's W26 hard-tech analysis